Individual Assignment 8

Group 3

11/3/2021

Exercises 8.4 Problem #8: In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative response variable. Now we will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable.

```
library(ISLR)
set.seed(2)
train = sample(1:nrow(Carseats),200)
Carseats.test = Carseats[-train,]
```

(d) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important.

```
#bagging
library(randomForest)
## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
set.seed(1)
bag.Carseats = randomForest(Sales~., data=Carseats, subset=train, mtry=10,
importance=TRUE)
bag.Carseats
##
## Call:
## randomForest(formula = Sales ~ ., data = Carseats, mtry = 10,
importance = TRUE, subset = train)
##
                  Type of random forest: regression
                        Number of trees: 500
##
## No. of variables tried at each split: 10
##
##
             Mean of squared residuals: 2.953352
                       % Var explained: 60.36
##
```

```
yhat.bag = predict(bag.Carseats, newdata = Carseats.test)
mean((yhat.bag-Carseats.test$Sales)^2)
## [1] 2.555523
#test MSE is 2.555523
importance(bag.Carseats)
                  %IncMSE IncNodePurity
##
## CompPrice
              25.64599599
                             213.096278
## Income
               5.67740403
                              76.385083
## Advertising 12.22093633
                             104.986108
## Population 0.59715138
                             61.294727
## Price
              56.50749020
                             535.237246
## ShelveLoc 41.05022443
                             283.698935
                            118.103039
## Age
              9.11485795
## Education -0.05758758
                             39.442533
## Urban
               0.79459812
                               8.828001
## US
               1.45384950
                               7.611851
#Price, ShelveLoc and CompPrice are most important
```

(e) Use random forests to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important. Describe the effect of m, the number of variables considered at each split, on the error rate obtained.

```
library(randomForest)
set.seed(1)
rf.Carseats = randomForest(Sales~.,data=Carseats, subset=train, mtry=3,
importance=TRUE)
rf.Carseats
##
## Call:
    randomForest(formula = Sales ~ ., data = Carseats, mtry = 3,
importance = TRUE, subset = train)
                  Type of random forest: regression
##
                        Number of trees: 500
## No. of variables tried at each split: 3
##
             Mean of squared residuals: 3.334325
##
##
                       % Var explained: 55.24
```

```
yhat.rf = predict(rf.Carseats, newdata = Carseats.test)
mean((yhat.rf-Carseats.test$Sales)^2)
## [1] 3.20635
#test MSE is 3.20635
importance(rf.Carseats)
##
                   %IncMSE IncNodePurity
## CompPrice
               11.34376552
                               158.83906
## Income
               6.40545807
                               121.23469
## Advertising 10.24191488
                               123.49672
## Population 0.09842459
                                98.17871
## Price
              36.27025901
                               403.38433
## ShelveLoc
              31.07819285
                               241.99669
## Age
               6.77865583
                               145.64539
## Education
               1.51515779
                                65.62168
## Urban
               1.14914879
                                14.60700
## US
                2.71845074
                                14.78064
#Price, ShelveLoc are most important
#A larger m will decrease test MSE
```

Problem #10: We now use boosting to predict Salary in the Hitters data set.

(a) Remove the observations for whom the salary information is unknown, and then log-transform the salaries.

```
Hitters=na.omit(Hitters)
log_salary=log(Hitters$Salary)
Hitters=data.frame(Hitters,log_salary)
```

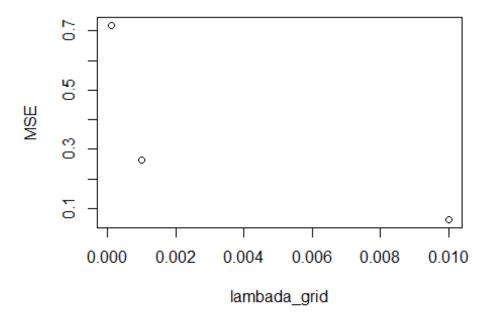
(b) Create a training set consisting of the first 200 observations, and a test set consisting of the remaining observations.

```
train=c(1:200)
test=-train
```

(c) Perform boosting on the training set with 1,000 trees for a range of values of the shrinkage parameter λ . Produce a plot with different shrinkage values on the x-axis and the corresponding training set MSE on the y-axis.

```
library(gbm)
## Loaded gbm 2.1.8
lambada_grid=c(0.01,0.001,0.0001)
MSE=rep(0,3)
for(i in 1:3){
   boost.Hitters= gbm(log_salary~.-Salary,Hitters[train,],distribution =
   "gaussian", n.trees=1000,interaction.depth=4,shrinkage = lambada_grid[i] )
yhat.boost=predict(boost.Hitters,newdata=Hitters[train,],n.trees=1000,interaction.depth=4)
```

```
tion.depth=4,shrinkage = lambada_grid[i] )
   MSE[i]=mean((yhat.boost - Hitters[train,]$log_salary)^2) }
plot(lambada_grid,MSE)
```

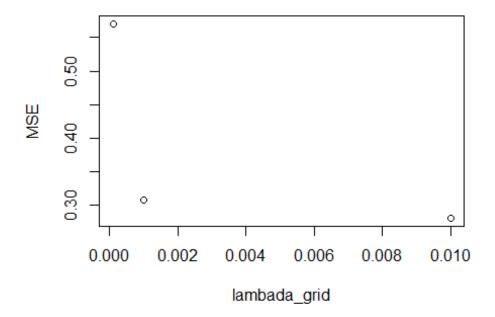


(d) Produce a plot with different shrinkage values on the x-axis and the corresponding test set MSE on the y-axis.

```
for(i in 1:3){
   boost.Hitters= gbm(log_salary~.-Salary,Hitters[train,],distribution =
   "gaussian", n.trees=1000,interaction.depth=4,shrinkage = lambada_grid[i] )

yhat.boost=predict(boost.Hitters,newdata=Hitters[test,],n.trees=1000,interact
ion.depth=4,shrinkage = lambada_grid[i] )
   MSE[i]=mean((yhat.boost - Hitters[test,]$log_salary)^2) }

plot(lambada_grid,MSE)
```



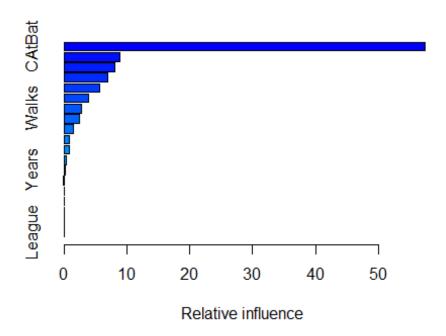
```
MSE
## [1] 0.2799206 0.3065863 0.5705474
```

(e) Compare the test MSE of boosting to the test MSE that results from applying two of the regression approaches seen in Chapters 3 and 6.

```
library(glmnet)
## 载入需要的程辑包: Matrix
## Loaded glmnet 4.1-2
x = model.matrix(log_salary~.-Salary,Hitters)[,-2]
y = Hitters$log_salary
grid = 10^seq(10, -2, length=100)
#ridge regression
ridge.mod = glmnet(x,y,alpha=0,lambda=grid,thresh = 1e-12)
set.seed(1)
cv.out = cv.glmnet(x[train,],y[train], alpha=0) ##default is 10-fold cross-
validation.
bestlam = cv.out$lambda.min
ridge.pred=predict(ridge.mod, s=bestlam, x=x[train,], y=y[train], newx=x[test,],e
xact = T)
mean((ridge.pred-y[test])^2)
## [1] 0.4521353
```

```
#LASSO
lasso.mod = glmnet(x,y,alpha=1,lambda=grid,thresh = 1e-12)
set.seed(1)
cv.out = cv.glmnet(x[train,],y[train], alpha=1) ##default is 10-fold cross-
validation.
bestlam = cv.out$lambda.min
lasso.pred=predict(lasso.mod, s=bestlam, x=x[train,], y=y[train], newx=x[test,],e
xact = T)
mean((lasso.pred-y[test])^2)
## [1] 0.4696766
#linear regression
glm.fit = glm(log_salary~.-Salary,data = Hitters ,subset = train)
yhat.glm = predict(glm.fit, newdata = Hitters[test,], type = "response")
mean((yhat.glm - Hitters[test,]$log_salary) ^ 2)
## [1] 0.4917959
# The test MSE of boosting With shrinkage of 0.01 is smaller than linear
regression, ridge regression and LASSO
```

(f) Which variables appear to be the most important predictors in the boosted model? summary(boost.Hitters)



```
## catBat catBat 57.436983418
```

```
## CRuns
                CRuns 8.843024418
## CHits
                CHits 8.051207483
## CWalks
               CWalks 6.904736267
## CRBI
                  CRBI 5.717567478
## AtBat
                AtBat 3.874379438
## Walks
                Walks 2.758656329
## CHmRun
               CHmRun 2,423226262
## Hits
                 Hits 1.541525651
## RBI
                  RBI 0.913012418
## Runs
                 Runs 0.801367327
## PutOuts
              PutOuts 0.382220942
                Years 0.116283565
## Years
## HmRun
                HmRun 0.104311732
## Assists
              Assists 0.072362761
## Errors
                Errors 0.027614873
## NewLeague NewLeague 0.011956870
## Division
             Division 0.011114482
               League 0.008448285
## League
```

#CAtBat appears to be the most important predictors in the boosted model

```
(g) Now apply bagging to the training set. What is the test set MSE for this approach?
set.seed(1)
bag.Hitters = randomForest(log_salary~.-Salary, data=Hitters, subset=train,
mtry=19, importance=TRUE)
bag.Hitters
##
## Call:
## randomForest(formula = log_salary ~ . - Salary, data = Hitters,
                                                                           mtry
= 19, importance = TRUE, subset = train)
                  Type of random forest: regression
##
                        Number of trees: 500
##
## No. of variables tried at each split: 19
##
##
             Mean of squared residuals: 0.2178554
##
                       % Var explained: 73.82
yhat.bag = predict(bag.Hitters, newdata = Hitters[test,])
mean((yhat.bag-Hitters[test,]$log_salary)^2)
## [1] 0.2301184
#test MSE is 0.2301184
```